

data-wrangling

October 22, 2019

```
<a href="https://cocl.us/corsera_da0101en_notebook_top">
  
```

Data Analysis with Python

Data Wrangling

Welcome!

By the end of this notebook, you will have learned the basics of Data Wrangling!

Table of content

```
<li><a href="#identify_handle_missing_values">Identify and handle missing values</a>
  <ul>
    <li><a href="#identify_missing_values">Identify missing values</a></li>
    <li><a href="#deal_missing_values">Deal with missing values</a></li>
    <li><a href="#correct_data_format">Correct data format</a></li>
  </ul>
</li>
<li><a href="#data_standardization">Data standardization</a></li>
<li><a href="#data_normalization">Data Normalization (centering/scaling)</a></li>
<li><a href="#binning">Binning</a></li>
<li><a href="#indicator">Indicator variable</a></li>
```

Estimated Time Needed: 30 min

What is the purpose of Data Wrangling?

Data Wrangling is the process of converting data from the initial format to a format that may be better for analysis.

What is the fuel consumption (L/100k) rate for the diesel car?

Import data

You can find the "Automobile Data Set" from the following link:
<https://archive.ics.uci.edu/ml/machine-learning-databases/autos/imports-85.data>. We will be using this data set throughout this course.

Import pandas

```
[2]: import pandas as pd
import matplotlib.pyplot as plt
```

Reading the data set from the URL and adding the related headers.

URL of the dataset

This dataset was hosted on IBM Cloud object click [HERE](#) for free storage

```
[3]: filename = "https://s3-api.us-geo.objectstorage.softlayer.net/cf-courses-data/
      ↪CognitiveClass/DA0101EN/auto.csv"
```

Python list headers containing name of headers

```
[4]: headers = ["symboling", "normalized-losses", "make", "fuel-type", "aspiration",
      ↪"num-of-doors", "body-style",
      ↪"drive-wheels", "engine-location", "wheel-base",
      ↪"length", "width", "height", "curb-weight", "engine-type",
      ↪"num-of-cylinders",
      ↪"engine-size", "fuel-system", "bore", "stroke", "compression-ratio", "horsepower",
      ↪"peak-rpm", "city-mpg", "highway-mpg", "price"]
```

Use the Pandas method `read_csv()` to load the data from the web address. Set the parameter "names" equal to the Python list "headers".

```
[5]: df = pd.read_csv(filename, names = headers)
```

Use the method `head()` to display the first five rows of the dataframe.

```
[6]: # To see what the data set looks like, we'll use the head() method.
      df.head()
```

```
[6]:  symboling normalized-losses      make fuel-type aspiration num-of-doors \
0         3             ?  alfa-romero    gas      std         two
1         3             ?  alfa-romero    gas      std         two
2         1             ?  alfa-romero    gas      std         two
3         2          164      audi      gas      std         four
4         2          164      audi      gas      std         four

      body-style drive-wheels engine-location  wheel-base  ...  engine-size  \
0  convertible      rwd      front      88.6  ...      130
1  convertible      rwd      front      88.6  ...      130
2   hatchback      rwd      front      94.5  ...      152
3      sedan      fwd      front      99.8  ...      109
4      sedan      4wd      front      99.4  ...      136

      fuel-system  bore  stroke  compression-ratio  horsepower  peak-rpm  city-mpg  \
0      mpfi  3.47   2.68           9.0          111      5000      21
1      mpfi  3.47   2.68           9.0          111      5000      21
2      mpfi  2.68   3.47           9.0          154      5000      19
3      mpfi  3.19   3.40          10.0          102      5500      24
4      mpfi  3.19   3.40           8.0          115      5500      18
```

	highway-mpg	price
0	27	13495
1	27	16500
2	26	16500
3	30	13950
4	22	17450

[5 rows x 26 columns]

As we can see, several question marks appeared in the dataframe; those are missing values which may hinder our further analysis.

So, how do we identify all those missing values and deal with them?

How to work with missing data?

Steps for working with missing data:

- identify missing data
- deal with missing data
- correct data format

Identify and handle missing values

Identify missing values

Convert "?" to NaN

In the car dataset, missing data comes with the question mark "?". We replace "?" with NaN (Not a Number), which is Python's default missing value marker, for reasons of computational speed and convenience. Here we use the function:

to replace A by B

```
[7]: import numpy as np

# replace "?" to NaN
df.replace("?", np.nan, inplace = True)
df.head(5)
```

```
[7]:   symboling normalized-losses      make fuel-type aspiration num-of-doors \
0         3             NaN  alfa-romero    gas      std         two
1         3             NaN  alfa-romero    gas      std         two
2         1             NaN  alfa-romero    gas      std         two
3         2            164      audi      gas      std         four
4         2            164      audi      gas      std         four

   body-style drive-wheels engine-location  wheel-base  ...  engine-size  \
0  convertible         rwd         front      88.6  ...        130
1  convertible         rwd         front      88.6  ...        130
2   hatchback         rwd         front      94.5  ...        152
```

3	sedan	fwd	front	99.8	...	109
4	sedan	4wd	front	99.4	...	136

	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm	city-mpg	\
0	mpfi	3.47	2.68	9.0	111	5000	21	
1	mpfi	3.47	2.68	9.0	111	5000	21	
2	mpfi	2.68	3.47	9.0	154	5000	19	
3	mpfi	3.19	3.40	10.0	102	5500	24	
4	mpfi	3.19	3.40	8.0	115	5500	18	

	highway-mpg	price
0	27	13495
1	27	16500
2	26	16500
3	30	13950
4	22	17450

[5 rows x 26 columns]

identify__missing__values

Evaluating for Missing Data

The missing values are converted to Python's default. We use Python's built-in functions to identify these missing values. There are two methods to detect missing data:

```
<li><b>.isnull()</b></li>
<li><b>.notnull()</b></li>
```

The output is a boolean value indicating whether the value that is passed into the argument is in fact missing data.

```
[8]: missing_data = df.isnull()
missing_data.head(5)
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	\
0	False	True	False	False	False	False	
1	False	True	False	False	False	False	
2	False	True	False	False	False	False	
3	False	False	False	False	False	False	
4	False	False	False	False	False	False	

	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	\
0	False	False	False	False	...	False	
1	False	False	False	False	...	False	
2	False	False	False	False	...	False	
3	False	False	False	False	...	False	
4	False	False	False	False	...	False	

	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm	\
0	False	False	False	False	False	False	
1	False	False	False	False	False	False	
2	False	False	False	False	False	False	
3	False	False	False	False	False	False	
4	False	False	False	False	False	False	

	city-mpg	highway-mpg	price
0	False	False	False
1	False	False	False
2	False	False	False
3	False	False	False
4	False	False	False

[5 rows x 26 columns]

"True" stands for missing value, while "False" stands for not missing value.

Count missing values in each column

Using a for loop in Python, we can quickly figure out the number of missing values in each column. As mentioned above, "True" represents a missing value, "False" means the value is present in the dataset. In the body of the for loop the method ".value_counts()" counts the number of "True" values.

```
[9]: for column in missing_data.columns.values.tolist():
      print(column)
      print(missing_data[column].value_counts())
      print("")
```

symboling

False 205

Name: symboling, dtype: int64

normalized-losses

False 164

True 41

Name: normalized-losses, dtype: int64

make

False 205

Name: make, dtype: int64

fuel-type

False 205

Name: fuel-type, dtype: int64

aspiration

False 205

Name: aspiration, dtype: int64

num-of-doors
False 203
True 2
Name: num-of-doors, dtype: int64

body-style
False 205
Name: body-style, dtype: int64

drive-wheels
False 205
Name: drive-wheels, dtype: int64

engine-location
False 205
Name: engine-location, dtype: int64

wheel-base
False 205
Name: wheel-base, dtype: int64

length
False 205
Name: length, dtype: int64

width
False 205
Name: width, dtype: int64

height
False 205
Name: height, dtype: int64

curb-weight
False 205
Name: curb-weight, dtype: int64

engine-type
False 205
Name: engine-type, dtype: int64

num-of-cylinders
False 205
Name: num-of-cylinders, dtype: int64

engine-size

```

False      205
Name: engine-size, dtype: int64

fuel-system
False      205
Name: fuel-system, dtype: int64

bore
False      201
True        4
Name: bore, dtype: int64

stroke
False      201
True        4
Name: stroke, dtype: int64

compression-ratio
False      205
Name: compression-ratio, dtype: int64

horsepower
False      203
True        2
Name: horsepower, dtype: int64

peak-rpm
False      203
True        2
Name: peak-rpm, dtype: int64

city-mpg
False      205
Name: city-mpg, dtype: int64

highway-mpg
False      205
Name: highway-mpg, dtype: int64

price
False      201
True        4
Name: price, dtype: int64

```

Based on the summary above, each column has 205 rows of data, seven columns containing missing data:

```
<li>"normalized-losses": 41 missing data</li>
```

- "num-of-doors": 2 missing data
- "bore": 4 missing data
- "stroke" : 4 missing data
- "horsepower": 2 missing data
- "peak-rpm": 2 missing data
- "price": 4 missing data

Deal with missing data

How to deal with missing data?

- drop data

 - a. drop the whole row

 - b. drop the whole column

- replace data

 - a. replace it by mean

 - b. replace it by frequency

 - c. replace it based on other functions

Whole columns should be dropped only if most entries in the column are empty. In our dataset, none of the columns are empty enough to drop entirely. We have some freedom in choosing which method to replace data; however, some methods may seem more reasonable than others. We will apply each method to many different columns:

Replace by mean:

- "normalized-losses": 41 missing data, replace them with mean
- "stroke": 4 missing data, replace them with mean
- "bore": 4 missing data, replace them with mean
- "horsepower": 2 missing data, replace them with mean
- "peak-rpm": 2 missing data, replace them with mean

Replace by frequency:

- "num-of-doors": 2 missing data, replace them with "four".
 -
 - Reason: 84% sedans is four doors. Since four doors is most frequent, it is most li

Drop the whole row:

- "price": 4 missing data, simply delete the whole row
 -
 - Reason: price is what we want to predict. Any data entry without price data cannot

Calculate the average of the column


```
[11]: avg_norm_loss = df["normalized-losses"].astype("float").mean(axis=0)
      print("Average of normalized-losses:", avg_norm_loss)
```

Average of normalized-losses: 122.0

Replace "NaN" by mean value in "normalized-losses" column

```
[12]: df["normalized-losses"].replace(np.nan, avg_norm_loss, inplace=True)
```

Calculate the mean value for 'bore' column

```
[13]: avg_bore=df['bore'].astype('float').mean(axis=0)
      print("Average of bore:", avg_bore)
```

Average of bore: 3.3297512437810943

Replace NaN by mean value

```
[ ]: df["bore"].replace(np.nan, avg_bore, inplace=True)
```

Question #1:

According to the example above, replace NaN in "stroke" column by mean.

```
[16]: # Write your code below and press Shift+Enter to execute
      stroke_means=df["stroke"].astype('float').mean(axis=0)
      print("Average of stroke:", stroke_means)
      df["stroke"].replace(np.nan,stroke_means, inplace = True)
```

Average of stroke: 3.2554228855721394

Double-click here for the solution.

Calculate the mean value for the 'horsepower' column:

```
[17]: avg_horsepower = df['horsepower'].astype('float').mean(axis=0)
      print("Average horsepower:", avg_horsepower)
```

Average horsepower: 104.25615763546799

Replace "NaN" by mean value:

```
[19]: df['horsepower'].replace(np.nan, avg_horsepower, inplace=True)
```

Calculate the mean value for 'peak-rpm' column:

```
[20]: avg_peakrpm=df['peak-rpm'].astype('float').mean(axis=0)
      print("Average peak rpm:", avg_peakrpm)
```

Average peak rpm: 5125.369458128079

Replace NaN by mean value:

```
[21]: df['peak-rpm'].replace(np.nan, avg_peakrpm, inplace=True)
```

To see which values are present in a particular column, we can use the ".value_counts()" method:

```
[22]: df['num-of-doors'].value_counts()
```

```
[22]: four      114
      two       89
      Name: num-of-doors, dtype: int64
```

We can see that four doors are the most common type. We can also use the ".idxmax()" method to calculate for us the most common type automatically:

```
[23]: df['num-of-doors'].value_counts().idxmax()
```

```
[23]: 'four'
```

The replacement procedure is very similar to what we have seen previously

```
[24]: #replace the missing 'num-of-doors' values by the most frequent
      df["num-of-doors"].replace(np.nan, "four", inplace=True)
```

Finally, let's drop all rows that do not have price data:

```
[25]: # simply drop whole row with NaN in "price" column
      df.dropna(subset=["price"], axis=0, inplace=True)

      # reset index, because we dropped two rows
      df.reset_index(drop=True, inplace=True)
```

```
[26]: df.head()
```

```
[26]:   symboling  normalized-losses      make fuel-type aspiration num-of-doors \
0         3         122  alfa-romero    gas      std         two
1         3         122  alfa-romero    gas      std         two
2         1         122  alfa-romero    gas      std         two
3         2         164      audi     gas      std         four
4         2         164      audi     gas      std         four

      body-style drive-wheels engine-location  wheel-base  ...  engine-size  \
0  convertible         rwd         front      88.6  ...      130
1  convertible         rwd         front      88.6  ...      130
2   hatchback         rwd         front      94.5  ...      152
3        sedan         fwd         front      99.8  ...      109
4        sedan         4wd         front      99.4  ...      136

      fuel-system  bore  stroke compression-ratio horsepower  peak-rpm city-mpg  \
0         mpfi  3.47   2.68              9.0          111     5000      21
```

1	mpfi	3.47	2.68	9.0	111	5000	21
2	mpfi	2.68	3.47	9.0	154	5000	19
3	mpfi	3.19	3.40	10.0	102	5500	24
4	mpfi	3.19	3.40	8.0	115	5500	18

	highway-mpg	price
0	27	13495
1	27	16500
2	26	16500
3	30	13950
4	22	17450

[5 rows x 26 columns]

Good! Now, we obtain the dataset with no missing values.

Correct data format

We are almost there!

The last step in data cleaning is checking and making sure that all data is in the correct format (int, float, text or other).

In Pandas, we use

.dtype() to check the data type

.astype() to change the data type

Lets list the data types for each column

[27]: `df.dtypes`

```
[27]: symboling          int64
normalized-losses      object
make                   object
fuel-type              object
aspiration             object
num-of-doors           object
body-style             object
drive-wheels           object
engine-location        object
wheel-base            float64
length                float64
width                  float64
height                 float64
curb-weight            int64
engine-type            object
num-of-cylinders       object
engine-size            int64
fuel-system            object
```

```

bore                object
stroke              object
compression-ratio   float64
horsepower          object
peak-rpm            object
city-mpg            int64
highway-mpg         int64
price               object
dtype: object

```

As we can see above, some columns are not of the correct data type. Numerical variables should have type 'float' or 'int', and variables with strings such as categories should have type 'object'. For example, 'bore' and 'stroke' variables are numerical values that describe the engines, so we should expect them to be of the type 'float' or 'int'; however, they are shown as type 'object'. We have to convert data types into a proper format for each column using the "astype()" method.

Convert data types to proper format

```

[28]: df[["bore", "stroke"]] = df[["bore", "stroke"]].astype("float")
      df[["normalized-losses"]] = df[["normalized-losses"]].astype("int")
      df[["price"]] = df[["price"]].astype("float")
      df[["peak-rpm"]] = df[["peak-rpm"]].astype("float")

```

Let us list the columns after the conversion

```

[29]: df.dtypes

```

```

[29]: symboling          int64
      normalized-losses  int64
      make               object
      fuel-type          object
      aspiration          object
      num-of-doors        object
      body-style          object
      drive-wheels        object
      engine-location      object
      wheel-base          float64
      length              float64
      width               float64
      height              float64
      curb-weight          int64
      engine-type          object
      num-of-cylinders      object
      engine-size          int64
      fuel-system          object
      bore                float64
      stroke              float64
      compression-ratio    float64

```

```
horsepower      object
peak-rpm        float64
city-mpg        int64
highway-mpg     int64
price           float64
dtype: object
```

Wonderful!

Now, we finally obtain the cleaned dataset with no missing values and all data in its proper format.

Data Standardization

Data is usually collected from different agencies with different formats. (Data Standardization is also a term for a particular type of data normalization, where we subtract the mean and divide by the standard deviation)

What is Standardization?

Standardization is the process of transforming data into a common format which allows the researcher to make the meaningful comparison.

Example

Transform mpg to L/100km:

In our dataset, the fuel consumption columns "city-mpg" and "highway-mpg" are represented by mpg (miles per gallon) unit. Assume we are developing an application in a country that accept the fuel consumption with L/100km standard

We will need to apply data transformation to transform mpg into L/100km?

The formula for unit conversion is

$$\text{L/100km} = 235 / \text{mpg}$$

We can do many mathematical operations directly in Pandas.

```
[30]: df.head()
```

```
[30]:   symboling  normalized-losses      make fuel-type aspiration \
0         3          122  alfa-romero    gas      std
1         3          122  alfa-romero    gas      std
2         1          122  alfa-romero    gas      std
3         2          164      audi    gas      std
4         2          164      audi    gas      std

   num-of-doors  body-style drive-wheels engine-location  wheel-base  ... \
0         two  convertible      rwd      front      88.6  ...
1         two  convertible      rwd      front      88.6  ...
2         two   hatchback      rwd      front      94.5  ...
3         four     sedan      fwd      front      99.8  ...
4         four     sedan      4wd      front      99.4  ...
```

	engine-size	fuel-system	bore	stroke	compression-ratio	horsepower	\
0	130	mpfi	3.47	2.68	9.0	111	
1	130	mpfi	3.47	2.68	9.0	111	
2	152	mpfi	2.68	3.47	9.0	154	
3	109	mpfi	3.19	3.40	10.0	102	
4	136	mpfi	3.19	3.40	8.0	115	

	peak-rpm	city-mpg	highway-mpg	price
0	5000.0	21	27	13495.0
1	5000.0	21	27	16500.0
2	5000.0	19	26	16500.0
3	5500.0	24	30	13950.0
4	5500.0	18	22	17450.0

[5 rows x 26 columns]

```
[36]: # Convert mpg to L/100km by mathematical operation (235 divided by mpg)
df['city-L/100km'] = 235/df["city-mpg"]

# check your transformed data
df.head()
```

	symboling	normalized-losses	make	fuel-type	aspiration	\
0	3	122	alfa-romero	gas	std	
1	3	122	alfa-romero	gas	std	
2	1	122	alfa-romero	gas	std	
3	2	164	audi	gas	std	
4	2	164	audi	gas	std	

	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	\
0	two	convertible	rwd	front	88.6	...	
1	two	convertible	rwd	front	88.6	...	
2	two	hatchback	rwd	front	94.5	...	
3	four	sedan	fwd	front	99.8	...	
4	four	sedan	4wd	front	99.4	...	

	bore	stroke	compression-ratio	horsepower	peak-rpm	city-mpg	highway-mpg	\
0	3.47	2.68	9.0	111	5000.0	21	27	
1	3.47	2.68	9.0	111	5000.0	21	27	
2	2.68	3.47	9.0	154	5000.0	19	26	
3	3.19	3.40	10.0	102	5500.0	24	30	
4	3.19	3.40	8.0	115	5500.0	18	22	

	price	city-L/100km	highway-L/100km
0	13495.0	11.190476	8.703704
1	16500.0	11.190476	8.703704

```

2  16500.0      12.368421      9.038462
3  13950.0       9.791667      7.833333
4  17450.0     13.055556     10.681818

```

[5 rows x 28 columns]

Question #2:

According to the example above, transform mpg to L/100km in the column of "highway-mpg", and change the name of column to "highway-L/100km".

```

[40]: # Write your code below and press Shift+Enter to execute
df["highway-L/100km"]=235/df["highway-mpg"]
df.rename(columns={"highway-mpg":'highway-L/100km'}, inplace=True)

#df.drop(['highway-mpg'],axis=1)

```

```

      □
↪ -----

      KeyError                                Traceback (most recent call
↪ last)

      ~/conda/envs/python/lib/python3.6/site-packages/pandas/core/indexes/base.
↪ py in get_loc(self, key, method, tolerance)
      2896             try:
      -> 2897                 return self._engine.get_loc(key)
      2898             except KeyError:

      pandas/_libs/index.pyx in pandas._libs.index.IndexEngine.get_loc()

      pandas/_libs/index.pyx in pandas._libs.index.IndexEngine.get_loc()

      pandas/_libs/index.pyx in pandas._libs.index.IndexEngine.
↪ _get_loc_duplicates()

      pandas/_libs/index.pyx in pandas._libs.index.IndexEngine.
↪ _maybe_get_bool_indexer()

      KeyError: 'highway-mpg'

```

During handling of the above exception, another exception occurred:

```
KeyError                                Traceback (most recent call
↳ last)

<ipython-input-40-3feeee4a5845a> in <module>
      1 # Write your code below and press Shift+Enter to execute
----> 2 df["highway-L/100km"]=235/df["highway-mpg"]
      3 df.rename(columns={"highway-mpg":'highway-L/100km'}, inplace=True)
      4
      5 #df.drop(['highway-mpg'],axis=1)

~/conda/envs/python/lib/python3.6/site-packages/pandas/core/frame.py in
↳ __getitem__(self, key)
    2978         if self.columns.nlevels > 1:
    2979             return self._getitem_multilevel(key)
-> 2980         indexer = self.columns.get_loc(key)
    2981         if is_integer(indexer):
    2982             indexer = [indexer]

~/conda/envs/python/lib/python3.6/site-packages/pandas/core/indexes/base.
↳ py in get_loc(self, key, method, tolerance)
    2897         return self._engine.get_loc(key)
    2898         except KeyError:
-> 2899         return self._engine.get_loc(self.
↳ _maybe_cast_indexer(key))
    2900         indexer = self.get_indexer([key], method=method,
↳ tolerance=tolerance)
    2901         if indexer.ndim > 1 or indexer.size > 1:

pandas/_libs/index.pyx in pandas._libs.index.IndexEngine.get_loc()

pandas/_libs/index.pyx in pandas._libs.index.IndexEngine.get_loc()

pandas/_libs/index.pyx in pandas._libs.index.IndexEngine.
↳ _get_loc_duplicates()

pandas/_libs/index.pyx in pandas._libs.index.IndexEngine.
↳ _maybe_get_bool_indexer()
```


KeyError: 'highway-mpg'

[Double-click here for the solution.](#)

Data Normalization

Why normalization?

Normalization is the process of transforming values of several variables into a similar range. Typical normalizations include scaling the variable so the variable average is 0, scaling the variable so the variance is 1, or scaling variable so the variable values range from 0 to 1

Example

To demonstrate normalization, let's say we want to scale the columns "length", "width" and "height"

Target: would like to Normalize those variables so their value ranges from 0 to 1.

Approach: replace original value by (original value)/(maximum value)

```
[41]: # replace (original value) by (original value)/(maximum value)
df['length'] = df['length']/df['length'].max()
df['width'] = df['width']/df['width'].max()
```

Question #3:

According to the example above, normalize the column "height".

```
[43]: # Write your code below and press Shift+Enter to execute
df['height'] = df['height']/df['height'].max()
```

[Double-click here for the solution.](#)

Here we can see, we've normalized "length", "width" and "height" in the range of [0,1].

Binning

Why binning?

Binning is a process of transforming continuous numerical variables into discrete categorical

Example:

In our dataset, "horsepower" is a real valued variable ranging from 48 to 288, it has 57 unique values. What if we only care about the price difference between cars with high horsepower, medium horsepower, and little horsepower (3 types)? Can we rearrange them into three 'bins' to simplify analysis?

We will use the Pandas method 'cut' to segment the 'horsepower' column into 3 bins

Example of Binning Data In Pandas

Convert data to correct format


```
[45]: df["horsepower"]=df["horsepower"].astype(int, copy=True)
```

Lets plot the histogram of horspower, to see what the distribution of horsepower looks like.

```
[48]: %matplotlib inline
import matplotlib as plt
from matplotlib import pyplot
plt.pyplot.hist(df["horsepower"])

# set x/y labels and plot title
plt.pyplot.xlabel("horsepower")
plt.pyplot.ylabel("count")
plt.pyplot.title("horsepower bins")
```

```
[48]: Text(0.5, 1.0, 'horsepower bins')
```



We would like 3 bins of equal size bandwidth so we use numpy's `linspace(start_value, end_value, numbers_generated)` function.

Since we want to include the minimum value of horsepower we want to set `start_value=min(df["horsepower"])`.

Since we want to include the maximum value of horsepower we want to set `end_value=max(df["horsepower"])`.

Since we are building 3 bins of equal length, there should be 4 dividers, so `numbers_generated=4`.

We build a bin array, with a minimum value to a maximum value, with bandwidth calculated above. The bins will be values used to determine when one bin ends and another begins.

```
[49]: bins = np.linspace(min(df["horsepower"]), max(df["horsepower"]), 4)
      bins
```

```
[49]: array([ 48.          , 119.33333333, 190.66666667, 262.          ])
```

We set group names:

```
[50]: group_names = ['Low', 'Medium', 'High']
```

We apply the function "cut" to determine what each value of "df['horsepower']" belongs to.

```
[52]: df['horsepower-binned'] = pd.cut(df['horsepower'], bins, labels=group_names,
      ↪include_lowest=True)
      df[['horsepower', 'horsepower-binned']].head(20)
```

```
[52]:
```

	horsepower	horsepower-binned
0	111	Low
1	111	Low
2	154	Medium
3	102	Low
4	115	Low
5	110	Low
6	110	Low
7	110	Low
8	140	Medium
9	101	Low
10	101	Low
11	121	Medium
12	121	Medium
13	121	Medium
14	182	Medium
15	182	Medium
16	182	Medium
17	48	Low
18	70	Low
19	70	Low

Lets see the number of vehicles in each bin.

```
[53]: df["horsepower-binned"].value_counts()
```


```
[53]: Low      153
      Medium   43
      High      5
      Name: horsepower-binned, dtype: int64
```

Lets plot the distribution of each bin.

```
[54]: %matplotlib inline
import matplotlib as plt
from matplotlib import pyplot
pyplot.bar(group_names, df["horsepower-binned"].value_counts())

# set x/y labels and plot title
plt.pyplot.xlabel("horsepower")
plt.pyplot.ylabel("count")
plt.pyplot.title("horsepower bins")
```

```
[54]: Text(0.5, 1.0, 'horsepower bins')
```



Check the dataframe above carefully, you will find the last column provides the bins for "horsepower"

We successfully narrow the intervals from 57 to 3!

Bins visualization

Normally, a histogram is used to visualize the distribution of bins we created above.

```
[55]: %matplotlib inline
import matplotlib as plt
from matplotlib import pyplot
```

```


a = (0,1,2)

# draw histogram of attribute "horsepower" with bins = 3
plt.pyplot.hist(df["horsepower"], bins = 3)

# set x/y labels and plot title
plt.pyplot.xlabel("horsepower")
plt.pyplot.ylabel("count")
plt.pyplot.title("horsepower bins")

```

[55]: Text(0.5, 1.0, 'horsepower bins')



The plot above shows the binning result for attribute "horsepower".

Indicator variable (or dummy variable)

What is an indicator variable?

An indicator variable (or dummy variable) is a numerical variable used to label categories. The

Why we use indicator variables?

So we can use categorical variables for regression analysis in the later modules.

Example

We see the column "fuel-type" has two unique values, "gas" or "diesel". Regression doesn't und

We will use the panda's method 'get_dummies' to assign numerical values to different categories

```
[56]: df.columns
```

```
[56]: Index(['symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration',  
          'num-of-doors', 'body-style', 'drive-wheels', 'engine-location',  
          'wheel-base', 'length', 'width', 'height', 'curb-weight', 'engine-type',  
          'num-of-cylinders', 'engine-size', 'fuel-system', 'bore', 'stroke',  
          'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg',  
          'highway-L/100km', 'price', 'city-L/100km', 'highway-L/100km',  
          'horsepower-binned'],  
          dtype='object')
```

get indicator variables and assign it to data frame "dummy_variable_1"

```
[62]: dummy_variable_1 = pd.get_dummies(df["fuel-type"])  
      dummy_variable_1.head(87)
```

```
[62]:
```

	diesel	gas
0	0	1
1	0	1
2	0	1
3	0	1
4	0	1
..
82	0	1
83	0	1
84	0	1
85	0	1
86	0	1

[87 rows x 2 columns]

change column names for clarity

```
[75]: dummy_variable_1.rename(columns={'fuel-type-diesel':'gas', 'fuel-type-gas':  
    ↪ 'diesel'}, inplace=True)  
      dummy_variable_1.head(15)
```

```
[75]:
```

	diesel	gas
0	0	1
1	0	1
2	0	1
3	0	1
4	0	1
5	0	1
6	0	1
7	0	1

8	0	1
9	0	1
10	0	1
11	0	1
12	0	1
13	0	1
14	0	1

We now have the value 0 to represent "gas" and 1 to represent "diesel" in the column "fuel-type". We will now insert this column back into our original dataset.

```
[76]: # merge data frame "df" and "dummy_variable_1"
df = pd.concat([df, dummy_variable_1], axis=1)

# drop original column "fuel-type" from "df"
df.drop("fuel-type", axis = 1, inplace=True)
```

```

      □
↳ -----

      KeyError                                Traceback (most recent call↳
↳ last)

      <ipython-input-76-a92dbd6eade8> in <module>
          3
          4 # drop original column "fuel-type" from "df"
----> 5 df.drop("fuel-type", axis = 1, inplace=True)

      ~/conda/envs/python/lib/python3.6/site-packages/pandas/core/frame.py in
↳ drop(self, labels, axis, index, columns, level, inplace, errors)
          4100         level=level,
          4101         inplace=inplace,
-> 4102         errors=errors,
          4103     )
          4104

      ~/conda/envs/python/lib/python3.6/site-packages/pandas/core/generic.py
↳ in drop(self, labels, axis, index, columns, level, inplace, errors)
          3912         for axis, labels in axes.items():
          3913             if labels is not None:
-> 3914                 obj = obj._drop_axis(labels, axis, level=level,
↳ errors=errors)
          3915
          3916         if inplace:
```

```
~/conda/envs/python/lib/python3.6/site-packages/pandas/core/generic.py
↳ in _drop_axis(self, labels, axis, level, errors)
    3963         labels_missing = (axis.get_indexer_for(labels) ==
↳ -1).any()
    3964         if errors == "raise" and labels_missing:
-> 3965             raise KeyError("{} not found in axis".
↳ format(labels))
    3966
    3967         slicer = [slice(None)] * self.ndim
```

KeyError: "['fuel-type'] not found in axis"

```
[74]: df.tail(15)
```

```
[74]:
```

	symboling	normalized-losses	make	aspiration	num-of-doors	\
186	3	256	volkswagen	1	two	
187	0	122	volkswagen	1	four	
188	0	122	volkswagen	0	four	
189	0	122	volkswagen	1	four	
190	-2	103	volvo	1	four	
191	-1	74	volvo	1	four	
192	-2	103	volvo	1	four	
193	-1	74	volvo	1	four	
194	-2	103	volvo	0	four	
195	-1	74	volvo	0	four	
196	-1	95	volvo	1	four	
197	-1	95	volvo	0	four	
198	-1	95	volvo	1	four	
199	-1	95	volvo	0	four	
200	-1	95	volvo	0	four	

	body-style	drive-wheels	engine-location	wheel-base	length	...	\
186	hatchback	fwd	front	94.5	0.796252	...	
187	sedan	fwd	front	100.4	0.865930	...	
188	sedan	fwd	front	100.4	0.865930	...	
189	wagon	fwd	front	100.4	0.879865	...	
190	sedan	rwd	front	104.3	0.907256	...	
191	wagon	rwd	front	104.3	0.907256	...	
192	sedan	rwd	front	104.3	0.907256	...	
193	wagon	rwd	front	104.3	0.907256	...	
194	sedan	rwd	front	104.3	0.907256	...	
195	wagon	rwd	front	104.3	0.907256	...	
196	sedan	rwd	front	109.1	0.907256	...	

197	sedan	rwd	front	109.1	0.907256	...
198	sedan	rwd	front	109.1	0.907256	...
199	sedan	rwd	front	109.1	0.907256	...
200	sedan	rwd	front	109.1	0.907256	...

	horsepower	peak-rpm	city-mpg	highway-L/100km	price	city-L/100km	\
186	90	5500.0	24	29	9980.0	9.791667	
187	110	5500.0	19	24	13295.0	12.368421	
188	68	4500.0	33	38	13845.0	7.121212	
189	88	5500.0	25	31	12290.0	9.400000	
190	114	5400.0	23	28	12940.0	10.217391	
191	114	5400.0	23	28	13415.0	10.217391	
192	114	5400.0	24	28	15985.0	9.791667	
193	114	5400.0	24	28	16515.0	9.791667	
194	162	5100.0	17	22	18420.0	13.823529	
195	162	5100.0	17	22	18950.0	13.823529	
196	114	5400.0	23	28	16845.0	10.217391	
197	160	5300.0	19	25	19045.0	12.368421	
198	134	5500.0	18	23	21485.0	13.055556	
199	106	4800.0	26	27	22470.0	9.038462	
200	114	5400.0	19	25	22625.0	12.368421	

	highway-L/100km	horsepower-binned	diesel	gas
186	8.103448	Low	0	1
187	9.791667	Low	0	1
188	6.184211	Low	1	0
189	7.580645	Low	0	1
190	8.392857	Low	0	1
191	8.392857	Low	0	1
192	8.392857	Low	0	1
193	8.392857	Low	0	1
194	10.681818	Medium	0	1
195	10.681818	Medium	0	1
196	8.392857	Low	0	1
197	9.400000	Medium	0	1
198	10.217391	Medium	0	1
199	8.703704	Low	1	0
200	9.400000	Low	0	1

[15 rows x 30 columns]

The last two columns are now the indicator variable representation of the fuel-type variable. It's all 0s and 1s now.

Question #4:

As above, create indicator variable to the column of "aspiration": "std" to 0, while "turbo" to 1.

```
[73]: # Write your code below and press Shift+Enter to execute
df_asp =pd.get_dummies(df['aspiration'])

df_asp.rename(columns = {'0':'std', '1':'turbo'})
```

```
[73]:      0  1
0      0  1
1      0  1
2      0  1
3      0  1
4      0  1
..    .. ..
196    0  1
197    1  0
198    0  1
199    1  0
200    1  0
```

[201 rows x 2 columns]

Double-click here for the solution.

Question #5:

Merge the new dataframe to the original dataframe then drop the column 'aspiration'

```
[ ]: # Write your code below and press Shift+Enter to execute
```

Double-click here for the solution.

save the new csv

```
[ ]: df.to_csv('clean_df.csv')
```

Thank you for completing this notebook

<p><img src="https://s3-api.us-geo.

About the Authors:

This notebook was written by Mahdi Noorian PhD, Joseph Santarcangelo, Bahare Talayian, Eric Xiao, Steven Dong, Parizad, Hima Vsudevan and Fiorella Wenver and Yi Yao.

Joseph Santarcangelo is a Data Scientist at IBM, and holds a PhD in Electrical Engineering. His research focused on using Machine Learning, Signal Processing, and Computer Vision to determine how videos impact human cognition. Joseph has been working for IBM since he completed his PhD.

Copyright © 2018 IBM Developer Skills Network. This notebook and its source code are released under the terms of the MIT License.